




Process

Jinwon Jeong

jin4812@korea.ac.kr

Reference

- The Linux Programming Interface
 - Published in October 2010
 - No Starch Press
 - ISBN 978-1-59327-220-3



Process

Overview of system calls

- fork(), exec(), wait(), and exit().
- Each of these system calls has variants.

#include <unistd.h>

#include <sys/types.h>

<https://www.guru99.com/c-gcc-install.html>

Overview of system calls

- **Fork()**

- Allows one process, the parent, to **create a new process**, the **child**.
- This is done by **making the new child process** an (almost) exact duplicate of the parent.
- The child obtains copies of the **parent's stack, data, heap, and text segments**.

Overview of system calls

- **Exit()**
 - **Terminates** a process.
 - Making all resources (memory, open file descriptors, and so on) used by the process available for subsequent reallocation by the kernel.
 - The ***status* argument is an integer** that determines the termination status for the process.
 - Using the *wait()* system call, the parent can **retrieve this status**.

Overview of system calls

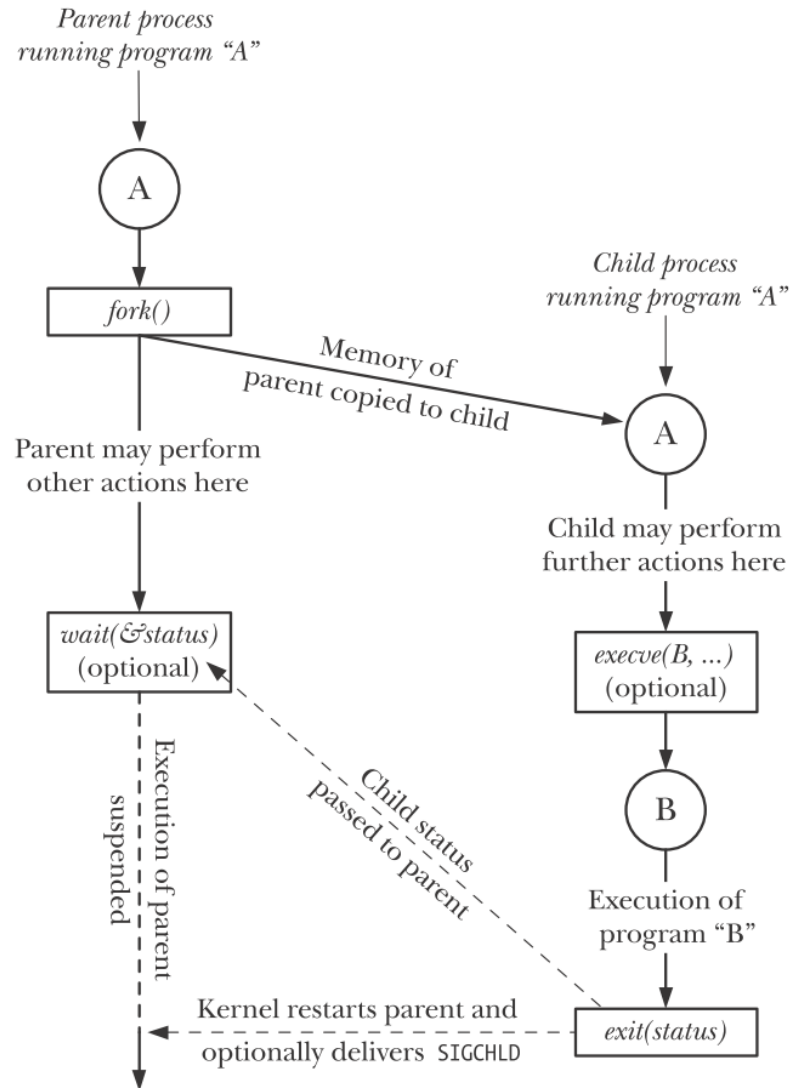
- **Wait()**

- (purpose) if a **child** of this process **has not yet terminated** by calling *exit()*, then *wait()* **suspends execution** of the process until one of its children has terminated.
- (purpose) the termination status of the child is returned in the status argument of *wait()*.
- `#include <wait.h>`

Overview of system calls

- **Exec()**
 - **Loads a new program** (*pathname*, with argument list *argv*, and environment list *envp*) into a process's memory.
 - The **existing program text is discarded**, and the stack, data, and heap segments are freshly created for the new program.
 - When Exec() invocation success, it doesn't return, but when fail, it returns -1.

Overview of system calls



Creating a New Process: *fork()*

- The `fork()` system call **creates a new process**, the ***child***, which is an almost exact duplicate of the calling process, the *parent*.

```
#include <unistd.h>
```

```
pid_t fork(void);
```

In parent: returns process ID of child on success, or -1 on error;
in successfully created child: always returns 0

Creating a New Process: *fork()*

- The key point to understanding *fork()* is to realize that after it has completed its work, **two processes exist**, and, in each process, **execution continues** from the point where *fork()* returns.
- The two processes are **executing the same program text**, but they have **separate copies** of the stack, data, and heap segments.

Creating a New Process: *fork()*

- The child's stack, data, and heap segments are initially exact **duplicates** of the corresponding parts **the parent's memory**.
- We can **distinguish** the two processes via the **value** returned from *fork()*.
- For the parent, *fork()* **returns the process ID** of the newly created child.
- For the child, ***fork()* returns 0**
 - If necessary, the child can obtain its own process ID using **getpid()**, and the process ID of its parent using **getppid()**.

Creating a New Process: *fork()*

- The following idiom is sometimes employed when calling *fork()*.

```
pid_t childPid;                /* Used in parent after successful fork()
                                to record PID of child */
switch (childPid = fork()) {
case -1:                        /* fork() failed */
    /* Handle error */

case 0:                         /* Child of successful fork() comes here */
    /* Perform actions specific to child */

default:                       /* Parent comes here after successful fork() */
    /* Perform actions specific to parent */
}
```

Creating a New Process: *fork()*

- It is important to realize that after a *fork()*, it is **indeterminate** which of the two processes is next scheduled to use the CPU.
 - In poorly written programs, this indeterminacy can lead to errors known as **race conditions**.

Creating a New Process: *fork()*

- **Copy on write(COW)**

- When parents want to make a child, copying all of parent's pages in memory makes the overhead.
- Just give a pointer that indicates the parent's pages to child.
- When child wants to modify the value inside the pages, COW gives a copy of that page to child for integrity.

Creating a New Process: *fork()*

fork1.c

```
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>
#include <stdlib.h>

static int idata = 111; /*Allocated in data segment*/

int main(int argc, char *argv[])
{
    int istack=222; /*Allocated in stack segment*/
    pid_t childPid;

    switch(childPid = fork())
    {
        case -1:
            exit(childPid);
        case 0:
            idata*=3;
            istack*=3;
            break;
        default:
            sleep(3);
            break;
    }
    printf("PID=%ld %s idata=%d istack=%d\n", (long)getpid(), (childPid==0) ?
    "(child)" : "(parent)", idata, istack);

    exit(childPid);
    return 0;
}
```


Creating a New Process: *fork()*

- The use of *sleep()* (in the code executed by the parent) in this program **permits the child to be scheduled** for the CPU **before the parent**.

```
root@kali:~# vi fork1.c
root@kali:~# gcc -o fork1 fork1.c
root@kali:~# ./fork1
PID=3876 (child) idata=333 istack=666
PID=3875 (parent) idata=111 istack=222
```

- The child process **gets its own copy** of the stack and data segments at the time of the *fork()*, and it is able to **modify variables** in these **segments without affecting the parent**.

Creating a New Process: *fork()*

fork2.c

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main(int argc, char *argv[])
{
    pid_t pid;
    int i=0;

    i++;
    printf("before(%d)\n", i);
    pid = fork();

    if(pid == 0)
        printf("child process(%d)\n", ++i);
    else if(pid>0)
        printf("parent(%d)\n", --i);
    else
        printf("fail");

    return 0;
}
```

Creating a New Process: *fork()*

fork3.c

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>
#include <stdlib.h>
#include <sys/wait.h>

#define max_pid 5
int i=5;
int main(){
    int j;
    //int status;
    pid_t pid[max_pid];

    for(j=0; j< max_pid; j++){
        pid[j] = fork();
        switch(pid[j]){
            case -1:
                perror("fork failed!\n");
                break;
            case 0 :
                i--;
                printf("child : %d, i = %d \n", getpid(), i);
                exit(2);
                break;
            default :
                i++;
                //wait(&status);
                printf("parent : %d, i = %d\n", getpid(), i);
        }
    }
    printf("last value i on parent's code : %d, i = %d \n",getpid(), i);
    return 0;
}
```

Executing a New Program: `exec()`

- The `exec()` system call **loads a new program** into a process's memory.
- During this operation, the **old program is discarded**, and the process's stack, data, and heap are **replaced** by those of the new program.

Executing a New Program: `execl()`

```
int execl(const char *pathname, const char *arg, ...  
          /* , (char *) NULL */);
```

None of the above returns on success; all return `-1` on error

- The **pathname** argument contains the pathname of the new program to be loaded into the process's memory.
- The **arg** argument specifies the command-line arguments to be passed to the new program.
- After an `execl()`, the **process ID** of the process remains the **same**, because the same process continues to exist.

Executing a New Program: `exec()`

- **The `exec()` Library Functions**

```
#include <unistd.h>

int execl(const char *pathname, const char *arg, ...
          /* , (char *) NULL, char *const envp[] */ );
int execlp(const char *filename, const char *arg, ...
          /* , (char *) NULL */);
int execvp(const char *filename, char *const argv[]);
int execv(const char *pathname, char *const argv[]);
int execl(const char *pathname, const char *arg, ...
          /* , (char *) NULL */);
```

None of the above returns on success; all return `-1` on error

Executing a New Program: execl()

execl1.c

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main(){
    printf("before executing ls -l\n");
    execl("/bin/ls", "ls", "-l", (char*)0);
    printf("where is my printf() result???\n");
    return 0;
}
```

Fork() + execl()

execl2.c

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main(){
    pid_t pid;
    printf("start\n");
    pid = fork();

    if(pid > 0){
        printf("parent\n");
        sleep(1);
    }else if(pid == 0){
        printf("child");
        execl("/bin/ls", "ls", "-l", (char*)0);
        printf("fail to execute ls -l \n");
    }else{
        printf("paraent fail to fork \n");
    }
    printf("bye\n");
    return 0;
}
```


Terminating a Process: *exit()*

- The *status* argument given to *exit()* defines the **termination status** of the process, which is available to the parent of this process when it calls *wait()*.

```
#include <unistd.h>

void exit(int status);
```

- By convention, a termination **status of 0** indicates that a process completed **successfully**, and a **nonzero** status value indicates that the process terminated **unsuccessfully**.

Waiting on a Child Process: *wait()*

- waits for **one of the children** of the calling process **to terminate** and
- returns the **termination status** of that child in the buffer pointed to by status.

```
#include <sys/wait.h>
```

```
pid_t wait(int *status);
```

Returns process ID of terminated child, or -1 on error

Waiting on a Child Process: *wait()*

wait1.c

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>
#include <stdlib.h>
#include <sys/wait.h>

int main(int argc, char *argv[])
{
    pid_t pid;
    int status;

    pid = fork();

    if(pid > 0)
    {
        /* parent process */
        printf("parent waiting..\n");
        wait(&status); // status = return value of child * 256
    }
    else if(pid == 0)
    {
        /* child process */
        sleep(1);
        printf("child: bye!\n");
        exit(2);
    }
    else
        printf("parent : fail to fork\n");

    printf("bye!\n");
    return 0;
}
```

Waiting on a Child Process: *wait()*

fork3.c

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>
#include <stdlib.h>
#include <sys/wait.h>

#define max_pid 5
int i=5;
int main(){
    int j;
    //int status;
    pid_t pid[max_pid];

    for(j=0; j< max_pid; j++){
        pid[j] = fork();
        switch(pid[j]){
            case -1:
                perror("fork failed!\n");
                break;
            case 0 :
                i--;
                printf("child : %d, i = %d \n", getpid(), i);
                exit(2);
                break;
            default :
                i++;
                //wait(&status);
                printf("parent : %d, i = %d\n", getpid(), i);
        }
    }
    printf("last value i on parent's code : %d, i = %d \n",getpid(), i);
    return 0;
}
```

Erase the annotation on fork3.c

Waiting on a Child Process: *waitpid()*

```
#include <sys/wait.h>

pid_t waitpid(pid_t pid, int *status_ptr, int options);
```

- Waits for specific child process to be end
- First : specific child process
- Second : child process's status
- Third : waitpid()'s option. 0, WNOHANG ...

Waiting on a Child Process: *waitpid()*

- WIFEXITED(status) : returns **true** if the child terminated normally
- WEXITSTATUS(status) : returns the **exit status** of the child.
- WTERMSIG(status) : returns the **number of the signal** that caused the child process to terminate.

Waiting on a Child Process: *waitpid()*

```
#include <stdio.h>
#include <unistd.h>
#include <wait.h>
#include <signal.h>

int main()
{
    int counter = 1;
    int status;
    pid_t pid1, pid2;
    pid_t pid_child;
    printf( "parent : Hello, my ID is %d\n", getpid());
    printf( "parent creates child 1.\n");
    pid1 = fork();

    switch(pid1){
        case -1 :
            printf( "fail, child 1\n");
            return -1;

        case 0 :
            printf( "child 1(%d) starts counting...!\n", getpid());
            while( 10 > counter ){
                printf( "child 1: %d\n", counter++);
                sleep(1);
            }
            return 1;

        default :

            printf( "parent creates child 2.\n");
            pid2 = fork();
            switch(pid2){
                case -1 :
                    printf( "fail, child 2\n");
                    return -1;

                case 0 :

                    printf( "child 2(%d) starts counting...!\n", getpid());
                    while( 100 > counter ){
                        printf( "child 2: %d\n", counter++);
                        sleep(1);
                    }
                    return 1;

                default :

                    printf( "parent : Im out of swith() now!! \n");
            }
    }
}
```

```
printf( "parent : I will start counting when child 1's counter value is 9...!\n");

pid_child = waitpid(pid1, &status, 0);

printf("parent : im start counting now... \n");
while(1){
    printf("parent : %d\n", counter++);
    sleep(1);
    if(counter == 8){
        break;
    }
}

printf( "terminated process is %d\n", pid_child);
if(WIFEXITED(status)){
    printf( "normal termination, status : %d\n", WEXITSTATUS(status));
}else{
    printf( "abnormal termination, status : %d\n", WIFSIGNALED(status));
}

kill(pid1, SIGKILL);
kill(pid2, SIGKILL);
return 0;
}
```

Waiting on a Child Process: *waitpid()*

```
#include <stdio.h>
#include <unistd.h>
#include <wait.h>
#include <signal.h>

int main()
{
    int counter = 1;
    int status;
    pid_t pid1, pid2;
    pid_t pid_child;
    printf( "parent : Hello, my ID is %d\n", getpid());
    printf( "parent creates child 1.\n");
    pid1 = fork();

    switch(pid1){
        case -1 :
            printf( "fail, child 1\n");
            return -1;

        case 0 :
            printf( "child 1(%d) starts counting...\n", getpid());
            while( 10 > counter ){
                printf( "child 1: %d\n", counter++);
                sleep(1);
            }
            return 1;

        default :

            printf( "parent creates child 2.\n");
            pid2 = fork();
            switch(pid2){
                case -1 :
                    printf( "fail, child 2\n");
                    return -1;

                case 0 :

                    printf( "child 2(%d) starts counting...\n", getpid());
                    while( 100 > counter ){
                        printf( "child 2: %d\n", counter++);
                        sleep(1);
                    }
                    return 1;

                default :

                    printf("parent : Im out of switch() now!! \n");
            }
    }

    printf( "parent : I will start counting now without block...\n");

    pid_child = waitpid(pid1, &status, WNOHANG);

    printf("parent : now im start counting!! \n");
    while(1){
        printf("parent : %d\n", counter++);
        sleep(1);
        if(counter == 8){
            break;
        }
    }

    printf( "terminated process is %d\n", pid_child);
    if(WIFEXITED(status)){
        printf( "normal termination, status : %d\n", WEXITSTATUS(status));
    }else{
        printf( "abnormal termination, status : %d\n", WIFSIGNALED(status));
    }
    kill(pid1, SIGKILL);
    kill(pid2, SIGKILL);
    return 0;
}
```


Assignment

Visit → <https://github.com/KU-OS/Process-Exercises>

Do : 01 ~ 05

KU-OS / Process-Exercises

<> Code ! Issues 🔗 Pull requests ▶ Actions 📁 Projects 📖 Wiki

🔗 main ▾ 🔗 1 branch 🏷 0 tags

root update	
📁 01	update
📁 02	update
📁 03	update
📁 04	update
📁 05	update
📄 LICENSE	Initial commit

Assignment

리눅스 환경(가상머신) 또는 macOS에서 실습 진행

1) 각 문제의 빈칸에 들어갈 코드 작성(전체 코드 작성할 필요 없음)

<P1/> XXX

<P2/> XXX

...

2) 각 문제마다 출력 결과 스크린샷(계정 보이게)

```
root@kali:~# gcc -o fork1 fork1.c
root@kali:~# ./fork1
PID=10878 (child) idata=333 istack=666
PID=10877 (parent) idata=111 istack=222
```

PDF파일로 제출