



Thread


Jinwon Jeong

jin4812@korea.ac.kr

Thread

Reference

- The Linux Programming Interface
 - Published in October 2010
 - No Starch Press
 - ISBN 978-1-59327-220-3



Threads: Introduction

- We describe POSIX threads, often known as *Pthreads*.
- Like processes, threads are a mechanism that permits an application to perform **multiple tasks concurrently**.
- A single process can **contain multiple threads**.
- All of these threads are **independently executing** the same program.
- They **all share** the same global memory and heap segments.

Threads: Introduction

- On a multiprocessor system, multiple threads can be **executed in parallel**.
 - If one thread is blocked on I/O, other threads are still eligible to execute.

Threads: Introduction

- Processes does have the following **limitations** in some applications:
 - It is **difficult to share information** between processes.
 - Since the **parent and child don't share memory** (other than the read-only text segment), we **must use some form** of interprocess communication in order to exchange information between processes.

Threads: Introduction

- Process creation with *fork()* is **relatively expensive**.
- Even with the copy-on-write technique, the need to duplicate various process attributes such as **page tables and file descriptor tables** means that a *fork()* call is **still time-consuming**.

Threads: Introduction

- Threads address both of these problems:
 - Sharing information between threads is easy and fast.
 - It is just a matter of copying data into shared (global or heap) variables
 - Synchronization techniques are needed.
- Thread creation is faster than process creation
 - Typically 10 times faster or better.
 - Thread creation is faster because many of the attributes that must be duplicated in a child created by fork() are instead shared between threads.

Threads: Introduction

- Pthreads data types

Data type	Description
<i>pthread_t</i>	Thread identifier
<i>pthread_mutex_t</i>	Mutex
<i>pthread_mutexattr_t</i>	Mutex attributes object
<i>pthread_cond_t</i>	Condition variable
<i>pthread_condattr_t</i>	Condition variable attributes object
<i>pthread_key_t</i>	Key for thread-specific data
<i>pthread_once_t</i>	One-time initialization control context
<i>pthread_attr_t</i>	Thread attributes object

- When compile, use *-pthread* or *-lpthread* option.

Threads: Introduction

- Thread Creation
 - The *pthread_create()* function creates a new thread.

```
#include <pthread.h>
```

```
int pthread_create(pthread_t *thread, const pthread_attr_t *attr,  
                  void *(*start)(void *), void *arg);
```

Return 0 on success, or a positive error number on error

Threads: Introduction

- After a call to *pthread_create()*, a program **has no guarantees** about which thread will next be scheduled to use the CPU.
- Programs that implicitly **rely on a particular order** of scheduling are open to the same sorts of **race conditions**.
- If we need to enforce a particular order of execution, we must use one of the **synchronization techniques**.

Threads: Introduction

- Thread Termination

- The execution of a thread terminates in one of the following ways:
 - The thread's start function performs a return specifying a return value for the thread.
 - The thread call *pthread_exit()*.
 - The thread is canceled using *pthread_cancel()*.
 - Any of the threads calls *exit()*, or the main thread performs a return.(in the *main()* function)
 - It causes all threads in the process to terminate immediately.

Threads: Introduction

- The *pthread_exit()* function terminates the calling thread, and specifies a return value that can be obtained in another thread by calling *pthread_join()*.

```
#include <pthread.h>
```

```
void pthread_exit(void *retval);
```

Threads: Introduction

```
1  #include<pthread.h>
2  #include<stdlib.h>
3  #include<stdio.h>
4
5  void *hello_thread (void *arg)
6  {
7      printf("thread : hello\n");
8      return arg;
9  }
10
11 main()
12 {
13     pthread_t tid;
14     int status;
15
16     status = pthread_create(&tid, NULL, hello_thread, NULL);
17     if(status != 0)
18         perror("create Thread");
19     pthread_exit(NULL);
20 }
```

Threads: Introduction

```
1 #include<stdlib.h>
2 #include<pthread.h>
3 #include<stdio.h>
4
5 #define NUM_THREADS 3
6
7 void *hello_thread(void *arg)
8 {
9     printf("Thread %ld : Hello, world\n", (long int)arg);
10    return arg;
11 }
12
13 main()
14 {
15     pthread_t tid [NUM_THREADS];
16     long int i, status;
17
18     for(i=0; i<NUM_THREADS; i++)
19     {
20         status = pthread_create(&tid[i], NULL, hello_thread, (void *)i);
21         if(status != 0)
22         {
23             fprintf(stderr, "create thread %ld : %ld\n", i, status);
24             exit(1);
25         }
26     }
27     pthread_exit(NULL);
28 }
```

Threads: Introduction

- Joining with a Terminated Thread.
 - The *pthread_join()* function **waits for the thread to terminate.**
 - If that thread has already terminated, *pthread_join()* returns immediately.

```
#include <pthread.h>
```

```
int pthread_join(pthread_t thread, void **retval);
```

Returns 0 on success, or a positive error number on error

Threads: Introduction

- If a thread is **not detaches**, then we **must join** with it using *pthread_join()*.
- If we fail to do this, then, when the thread terminates, it produces the thread equivalent of a **zombie process**.
- Aside from **wasting system resources**, we **won't be able to create** additional threads.

Threads: Introduction

join.c

```
#include <pthread.h>
#include <stdio.h>
#include <unistd.h>
#include <stdlib.h>
|
#define NUM_THREADS 5
pthread_t tid[NUM_THREADS];

void *hello_thread(void *arg){

    int i = 0;
    while(1){
        if(i == 5){
            break;
        }
        i++;
        sleep(1);
    }
    printf("thread%ld terminated... \n", (intptr_t) arg);
    pthread_exit( (void*)(intptr_t)i);
}
```

Threads: Introduction

```
void main(){

    long int i, status;
    printf("pid : %d\n", getpid());

    for(i=0; i<=NUM_THREADS; i++){
        status = pthread_create(&tid[i], NULL, &hello_thread, (void *)(&i));
        if(status != 0){
            fprintf(stderr, "create thread %d, status : %d\n", i, status);
            exit(1);
        }
        printf("thread%d is created...\n", i);
    }

    for(i=0; i<=NUM_THREADS; i++){
        void *return_value;
        pthread_join(tid[i], &return_value);
        printf("thread%d's return value is %d\n", i, (intptr_t)return_value);
    }
}
```

Threads: Introduction

- Detaching a Thread
 - By default, a thread is *joinable*, meaning that when it terminates, another thread can obtain its return status using *pthread_join()*.
 - Sometimes we don't care about the thread's return status.
 - We simply want the system to automatically clean up and remove the thread when it terminates.

Threads: Introduction

- In this case, we can mark the thread as detached, by making a call to `pthread_detach()` specifying the thread's identifier in `thread`.

```
#include <pthread.h>
```

```
int pthread_detach(pthread_t thread);
```

Returns 0 on success, or a positive error number on error

Threads: Introduction

- Once a thread has been detached, it is no longer possible to use *pthread_join()*.
 - The thread can't be made joinable again.

Threads: Introduction

detach.c

```
#include <pthread.h>
#include <stdio.h>
#include <unistd.h>
#include <stdlib.h>

#define NUM_THREADS 5

int end[NUM_THREADS];
pthread_t tid[NUM_THREADS];

void *hello_thread(void *arg){

    int i;
    double result = 0.0;
    int state;
    state = pthread_detach(pthread_self());
    if(state == 0){
        printf("thread%d(%d) is detached!\n", (int)(intptr_t) arg, getpid());
    }else{
        printf("pthread_detach() got a error.,,\n");
    }

    while(!end[(int)(intptr_t) arg]){
        for(i=0; i<50000; i++){
            result = result + (double)random();
        }
    }
    printf("thread%d's result : %e\n", (int)(intptr_t) arg, result);
    printf("thread%d terminated...\n", (int)(intptr_t) arg);
}
```

Threads: Introduction

```
void main(){

    long int i, status;
    printf("pid : %d\n", getpid());
    for(i=0;i<=NUM_THREADS; i++){
        end[i] = 0;
        status = pthread_create(&tid[i], NULL, hello_thread, (void *)(&i));
        if(status !=0){
            fprintf(stderr, "create thread %d, status : %d\n", i, status);
            exit(1);
        }
        printf("thread%d is created...\n", i);
    }
    sleep(1);
    for(i=0;i<=NUM_THREADS; i++){
        end[i] = 1;
    }
    sleep(5);
    printf("++pid%d is finished...++\n", getpid());

    return;

}
```


Threads: Thread Synchronization

- We describe two tools that threads can use to synchronize their actions:
 - Mutexes
 - Condition variables

Threads: Thread Synchronization

- **Mutexes** allow threads to synchronize their use of a shared resource.
 - So that one thread doesn't try to access a shared variable at the same time as another thread is modifying it.
- **Condition variables** perform a complementary task.
 - They allow threads to inform each other that a shared variable (or other shared resource) has changed state.

Threads: Thread Synchronization

- The term **critical section** is used to refer to **a section of code** that accesses a shared resource and whose execution should be *atomic* (*All or Nothing*).
 - Its execution **should not be interrupted** by another thread that simultaneously accesses **the same shared resource**.

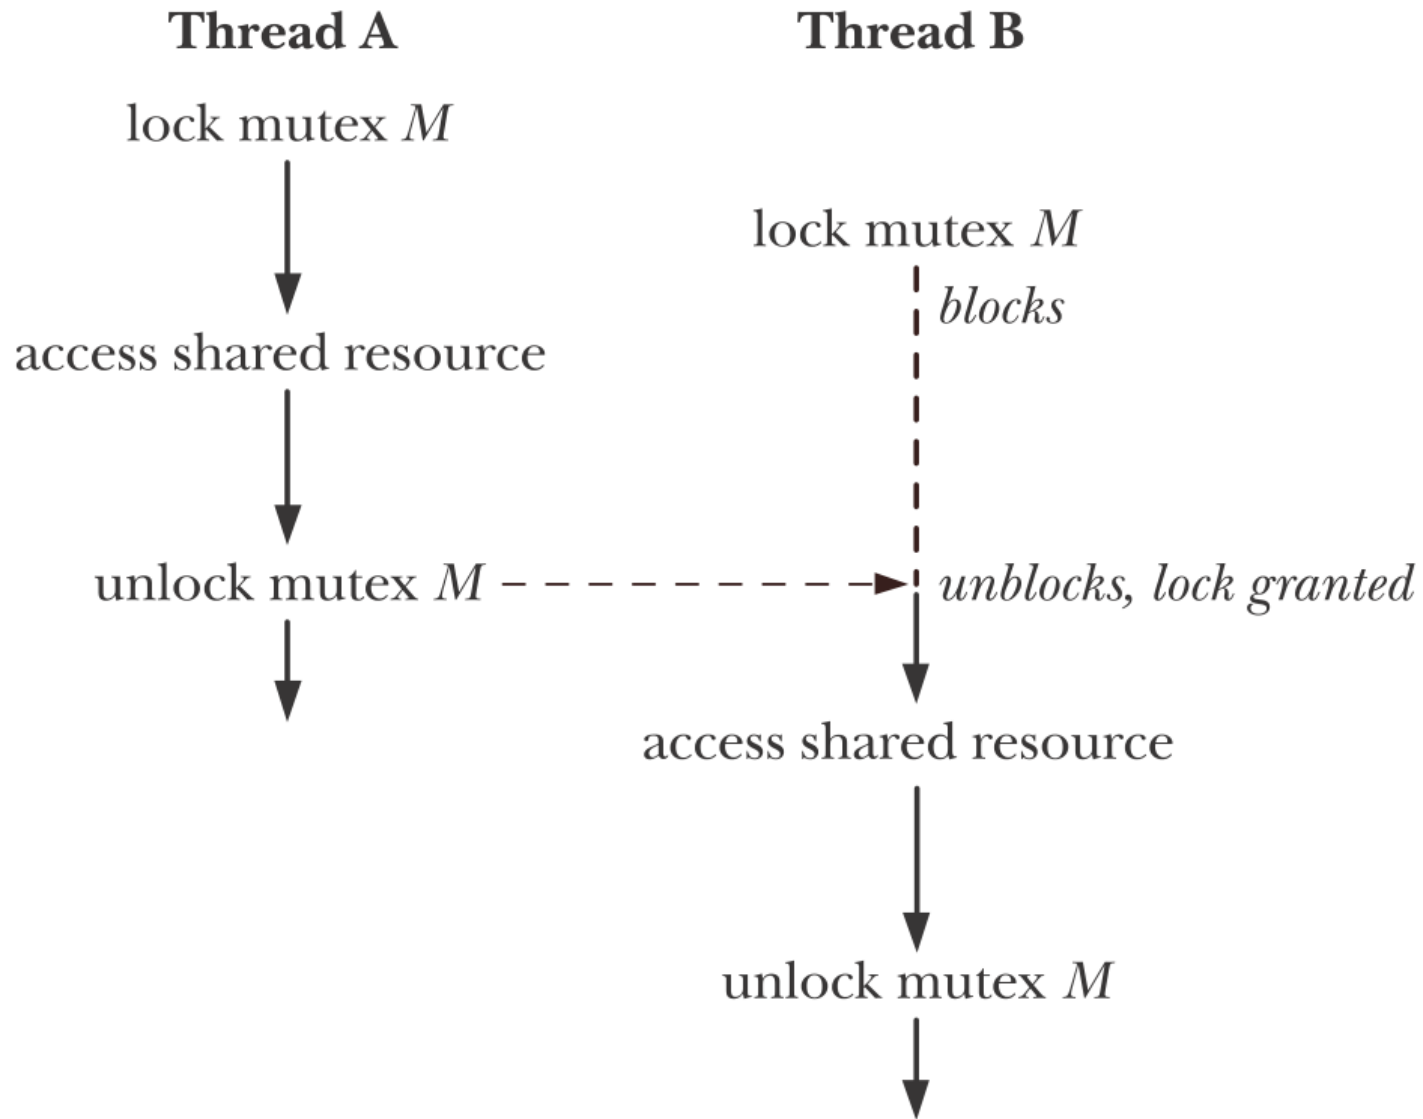
Mutexes

- To avoid the problems that can occur when threads try to **update a shared variable**, we must use a *mutex* (short for *mutual exclusion*) to ensure that **only one thread at a time** can access the variable.
- A mutex has two states:
 - *locked*
 - *unlocked*

Mutexes

- At any moment, **at most one thread** may hold the lock on a mutex.
- Attempting to lock a mutex that is **already locked** either **blocks or fails** with an error.

Mutexes



Mutexes

- A mutex is a variable of the type *pthread_mutex_t*.
- Before it can be used, a mutex must always be **initialized**.

```
#include <pthread.h>
```

```
int pthread_mutex_init(pthread_mutex_t *mutex, const pthread_mutexattr_t *attr);
```

Returns 0 on success, or a positive error number on error

```
#include <pthread.h>
```

```
int pthread_mutex_destroy(pthread_mutex_t *mutex);
```

Returns 0 on success, or a positive error number on error

Mutexes

– Locking and Unlocking a Mutex.

- After initialization, a **mutex is unlocked**.
- To lock and unlock a mutex, we use the *pthread_mutex_lock()* and *pthread_mutex_unlock()* functions.

```
#include <pthread.h>
```

```
int pthread_mutex_lock(pthread_mutex_t *mutex);
```

```
int pthread_mutex_unlock(pthread_mutex_t *mutex);
```

Both return 0 on success, or a positive error number on error

Mutexes

```
1 #include <stdlib.h>
2 #include <pthread.h>
3 #include <stdio.h>
4
5 #define NUM_THREADS 3
6
7 pthread_mutex_t mutex;
8 long int sum;
9 void *mutex_thread(void *arg)
10 {
11     pthread_mutex_lock (&mutex);
12     sum +=(long int) arg;
13     pthread_mutex_unlock (&mutex);
14
15     return arg;
16 }
17
```

```
18 int main (int argc, char *argv[])
19 {
20     pthread_t tid[NUM_THREADS];
21     long int arg[NUM_THREADS], i;
22     void *result;
23     int status;
24
25     if (argc < 4)
26     {
27         fprintf( stderr, "usage: mutexthread number1 number2 number3 \n");
28         exit(1);
29     }
30
31     for (i=0; i< NUM_THREADS; i++)
32         arg[i] = atoi (argv[i+1]);
33
34     pthread_mutex_init (&mutex, NULL);
35
36     for (i=0; i< NUM_THREADS; i++)
37     {
38         status = pthread_create ( &tid[i], NULL, mutex_thread, (void *) arg [i]);
39         if (status !=0)
40         {
41             fprintf(stderr, "create thread %d: %d" , (int)i, (int)status);
42             exit(1);
43         }
44     }
45     for (i=0; i< NUM_THREADS; i++)
46     {
47         status = pthread_join ( tid[i], &result);
48         if (status !=0)
49         {
50             fprintf(stderr, "join thread %d: %d" , (int)i, (int)status);
51             exit(1);
52         }
53     }
54
55     status = pthread_mutex_destroy (&mutex);
56     if (status != 0)
57         perror ("Desotry mutex");
58     printf(" sum is %ld\n", sum);
59     pthread_exit (result);
60 }
```

Assignment

Visit → <https://github.com/KU-OS/Threads-Exercises>

Do : 01 ~ 07

KU-OS / Threads-Exercises

<> Code ! Issues 🔗 Pull requests ▶ Actions 📁 Projects 📖 Wiki

🔗 main ▾ 🔗 1 branch 🔖 0 tags

root update		
📁	01	update
📁	02	update
📁	03	update
📁	04	update
📁	05	update
📁	06	update
📁	07	update
📄	LICENSE	Initial commit

Assignment

리눅스 환경(가상머신) 또는 macOS에서 실습 진행

1) 각 문제의 빈칸에 들어갈 코드 작성(전체 코드 작성할 필요 없음)

<P1/> XXX

<P2/> XXX

...

! 06과 07번 문제는 빈칸에 왜 그러한 코드를 작성하였는지에 대한 이유를 전체 코드의 흐름과 연관하여 서술할 것.

[07번 문제는 빈칸이 없으므로 적절한 공간에 코드를 작성하고 이유 서술.]

2) 각 문제마다 출력 결과 스크린샷(계정 보이게)

```
root@kali:~# gcc -o thread1 thread1.c -lpthread
root@kali:~# ./thread1
thread : hello
```

PDF 파일로 제출